

✓ Name :- Dev Sharma

CU24250269

BTech CSE 3rd Year SECTION - "B"

Roll No. :- 17

=====

Experiment No. 1

Experiment Name

Real-World Data Collection, Cleaning and Preprocessing

using Python

=====

Aim:

To collect a real-world dataset from publicly available

sources and perform data preprocessing by handling missing

values, duplicate records, inconsistent formats, categorical

variables, and outliers using Python.

GITHUB LINK :-[https://github.com/DevSharma1912/Mastery-in-Data-Analytics-and-Visualizations-and-Career-Advancement-Technical/tree/main/Mastery%20in%20Data%20Analytics%20and%20Visualizations%20and%20Career%20Advancement%20\(Technical\)/Experiment1](https://github.com/DevSharma1912/Mastery-in-Data-Analytics-and-Visualizations-and-Career-Advancement-Technical/tree/main/Mastery%20in%20Data%20Analytics%20and%20Visualizations%20and%20Career%20Advancement%20(Technical)/Experiment1)

Step 1: Import Required Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from tkinter import Tk, filedialog
from sklearn.preprocessing import LabelEncoder, StandardScaler
```

Step 2: Upload / Select Dataset from Computer

```
import io
from google.colab import files

print("Please select a CSV file to upload.")

uploaded = files.upload()

if not uploaded:
    print("No dataset selected.")
else:
    file_name = next(iter(uploaded)) # Get the name of the first uploaded file
    df = pd.read_csv(io.BytesIO(uploaded[file_name]))

    print("Dataset loaded successfully.")
    print("Selected Dataset:", file_name)
```

Please select a CSV file to upload.
Choose Files No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving WA_Fn-UseC_-HR-Employee-Attrition.csv to WA_Fn-UseC_-HR-Employee-Attrition.csv
Dataset loaded successfully.
Selected Dataset: WA_Fn-UseC_-HR-Employee-Attrition.csv

Step 3: Display First Five Records

```
print("\n===== FIRST FIVE RECORDS =====")

print(df.head())
```

```
===== FIRST FIVE RECORDS =====
   Age  Attrition  BusinessTravel  DailyRate  Department  \
0   41         Yes      Travel_Rarely    1102         Sales
1   49         No    Travel_Frequently    279  Research & Development
2   37         Yes      Travel_Rarely   1373  Research & Development
3   33         No    Travel_Frequently   1392  Research & Development
4   27         No      Travel_Rarely    591  Research & Development

   DistanceFromHome  Education  EducationField  EmployeeCount  EmployeeNumber  \
0                1         2  Life Sciences                1                1
1                8         1  Life Sciences                1                2
2                2         2         Other                1                4
3                3         4  Life Sciences                1                5
4                2         1         Medical                1                7

   ...  RelationshipSatisfaction  StandardHours  StockOptionLevel  \
0   ...                        1                80                0
1   ...                        4                80                1
2   ...                        2                80                0
3   ...                        3                80                0
4   ...                        4                80                1

   TotalWorkingYears  TrainingTimesLastYear  WorkLifeBalance  YearsAtCompany  \
0                8                0                1                6
1               10                3                3               10
2                7                3                3                0
3                8                3                3                8
4                6                3                3                2

   YearsInCurrentRole  YearsSinceLastPromotion  YearsWithCurrManager
0                4                0                5
1                7                1                7
2                0                0                0
3                7                3                0
4                2                2                2
```

[5 rows x 35 columns]

Step 4: Display Dataset Information

```
print("\n===== DATASET INFORMATION =====")
```

```
print(df.info())
```

```
===== DATASET INFORMATION =====
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Age                   1470 non-null  int64
 1   Attrition             1470 non-null  object
 2   BusinessTravel        1470 non-null  object
 3   DailyRate             1470 non-null  int64
 4   Department            1470 non-null  object
 5   DistanceFromHome      1470 non-null  int64
 6   Education             1470 non-null  int64
 7   EducationField        1470 non-null  object
 8   EmployeeCount         1470 non-null  int64
 9   EmployeeNumber        1470 non-null  int64
10   EnvironmentSatisfaction 1470 non-null  int64
11   Gender                1470 non-null  object
12   HourlyRate            1470 non-null  int64
13   JobInvolvement        1470 non-null  int64
14   JobLevel              1470 non-null  int64
15   JobRole               1470 non-null  object
16   JobSatisfaction       1470 non-null  int64
17   MaritalStatus         1470 non-null  object
18   MonthlyIncome         1470 non-null  int64
19   MonthlyRate           1470 non-null  int64
20   NumCompaniesWorked    1470 non-null  int64
21   Over18                1470 non-null  object
22   OverTime              1470 non-null  object
23   PercentSalaryHike     1470 non-null  int64
24   PerformanceRating     1470 non-null  int64
25   RelationshipSatisfaction 1470 non-null  int64
26   StandardHours         1470 non-null  int64
27   StockOptionLevel      1470 non-null  int64
28   TotalWorkingYears     1470 non-null  int64
29   TrainingTimesLastYear 1470 non-null  int64
30   WorkLifeBalance       1470 non-null  int64
31   YearsAtCompany        1470 non-null  int64
32   YearsInCurrentRole    1470 non-null  int64
33   YearsSinceLastPromotion 1470 non-null  int64
34   YearsWithCurrManager  1470 non-null  int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
None
```

Step 5: Display Statistical Description

```
print("\n===== STATISTICAL DESCRIPTION =====")
```

```
print(df.describe(include="all"))
```

```
===== STATISTICAL DESCRIPTION =====
count    1470.000000    Age  Attrition  BusinessTravel    DailyRate  \
unique           NaN                1470          1470    1470.000000
top            NaN                No          Travel_Rarely      NaN
freq           NaN                1233          1043      NaN
mean        36.923810      NaN                NaN    802.485714
std          9.135373      NaN                NaN    403.509100
min         18.000000      NaN                NaN    102.000000
25%         30.000000      NaN                NaN    465.000000
50%         36.000000      NaN                NaN    802.000000
75%         43.000000      NaN                NaN   1157.000000
max         60.000000      NaN                NaN   1499.000000

count    Department  DistanceFromHome  Education  EducationField  \
unique           3                NaN                NaN                6
top  Research & Development      NaN                NaN    Life Sciences
freq           961                NaN                NaN                606
mean           NaN          9.192517          2.912925                NaN
std           NaN          8.106864          1.024165                NaN
min           NaN          1.000000          1.000000                NaN
25%           NaN          2.000000          2.000000                NaN
50%           NaN          7.000000          3.000000                NaN
75%           NaN          14.000000          4.000000                NaN
max           NaN          29.000000          5.000000                NaN

count    EmployeeCount  EmployeeNumber  ...  RelationshipSatisfaction  \
unique           NaN                NaN                NaN                NaN
top            NaN                NaN                NaN                NaN
freq           NaN                NaN                NaN                NaN
mean           1.0          1024.865306  ...          2.712245
std           0.0          602.024335  ...          1.081209
min           1.0          1.000000  ...          1.000000
25%           1.0          491.250000  ...          2.000000
50%           1.0          1020.500000  ...          3.000000
75%           1.0          1555.750000  ...          4.000000
max           1.0          2068.000000  ...          4.000000

count    StandardHours  StockOptionLevel  TotalWorkingYears  \
unique           NaN                NaN                NaN
top            NaN                NaN                NaN
freq           NaN                NaN                NaN
mean           80.0          0.793878          11.279592
std           0.0          0.852077          7.780782
min           80.0          0.000000          0.000000
25%           80.0          0.000000          6.000000
50%           80.0          1.000000          10.000000
75%           80.0          1.000000          15.000000
max           80.0          3.000000          40.000000

count    TrainingTimesLastYear  WorkLifeBalance  YearsAtCompany  \
unique           NaN                NaN                NaN
top            NaN                NaN                NaN
```

Step 6: Check Dataset Shape

```
print("\n===== DATASET SHAPE =====")
```

```
print("Number of Rows:", df.shape[0])
print("Number of Columns:", df.shape[1])
```

```
===== DATASET SHAPE =====
Number of Rows: 1470
Number of Columns: 35
```

Step 7: Identify Missing Values

```
print("\n===== MISSING VALUES =====")

missing_values = df.isnull().sum()

print(missing_values)
```

```
===== MISSING VALUES =====
Age                                0
Attrition                         0
BusinessTravel                    0
DailyRate                        0
Department                       0
DistanceFromHome                 0
Education                        0
EducationField                    0
EmployeeCount                     0
EmployeeNumber                   0
EnvironmentSatisfaction           0
Gender                           0
HourlyRate                       0
JobInvolvement                   0
JobLevel                         0
JobRole                          0
JobSatisfaction                   0
MaritalStatus                    0
MonthlyIncome                    0
MonthlyRate                      0
NumCompaniesWorked               0
Over18                           0
OverTime                         0
PercentSalaryHike                0
PerformanceRating                0
RelationshipSatisfaction          0
StandardHours                    0
StockOptionLevel                 0
TotalWorkingYears                0
TrainingTimesLastYear            0
WorkLifeBalance                  0
YearsAtCompany                   0
YearsInCurrentRole                0
YearsSinceLastPromotion          0
YearsWithCurrManager             0
dtype: int64
```

Step 8: Handle Missing Values

```
# Numerical columns -> Median
numerical_columns = df.select_dtypes(
    include=np.number
).columns

for column in numerical_columns:
    if df[column].isnull().sum() > 0:
        df[column] = df[column].fillna(
            df[column].median()
        )

# Categorical columns -> Mode
categorical_columns = df.select_dtypes(
    exclude=np.number
).columns

for column in categorical_columns:
    if df[column].isnull().sum() > 0:
        mode_value = df[column].mode()

        if len(mode_value) > 0:
            df[column] = df[column].fillna(
                mode_value[0]
            )

print("\n===== MISSING VALUE HANDLING =====")

print("Numerical missing values filled using Median.")
print("Categorical missing values filled using Mode.")

print("\nRemaining Missing Values:")

print(df.isnull().sum().sum())
```

```
===== MISSING VALUE HANDLING =====
Numerical missing values filled using Median.
Categorical missing values filled using Mode.

Remaining Missing Values:
0
```

Step 9: Detect Duplicate Records

```
duplicate_count = df.duplicated().sum()

print("\n===== DUPLICATE RECORDS =====")

print("Number of Duplicate Records:", duplicate_count)

# Remove duplicates
df = df.drop_duplicates()

print("Duplicate Records Removed Successfully.")
print("Remaining Rows:", len(df))
```

```
===== DUPLICATE RECORDS =====
Number of Duplicate Records: 0
Duplicate Records Removed Successfully.
Remaining Rows: 1470
```

Step 10: Standardize Categorical Values

```
print("\n===== CATEGORICAL DATA =====")

categorical_columns = df.select_dtypes(
    exclude=np.number
).columns

for column in categorical_columns:

    # Remove unnecessary spaces
    df[column] = df[column].astype(str).str.strip()

    # Convert text values to consistent lowercase format
    df[column] = df[column].str.lower()

print("Categorical values standardized successfully.")
```

```
===== CATEGORICAL DATA =====
Categorical values standardized successfully.
```

Step 11: Encode Categorical Variables

```
encoded_df = df.copy()

categorical_columns = encoded_df.select_dtypes(
    exclude=np.number
).columns

label_encoder = LabelEncoder()

for column in categorical_columns:

    encoded_df[column] = label_encoder.fit_transform(
        encoded_df[column].astype(str)
    )

print("Categorical variables encoded successfully.")
print("Encoding Method: Label Encoding")
```

```
Categorical variables encoded successfully.
Encoding Method: Label Encoding
```

Step 12: Detect Outliers using IQR Method

```
print("\n===== OUTLIER DETECTION =====")

numeric_columns = encoded_df.select_dtypes(
    include=np.number
).columns

outlier_summary = {}

for column in numeric_columns:

    Q1 = encoded_df[column].quantile(0.25)
    Q3 = encoded_df[column].quantile(0.75)

    IQR = Q3 - Q1

    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    outliers = encoded_df[
        (encoded_df[column] < lower_bound) |
        (encoded_df[column] > upper_bound)
    ]

    outlier_summary[column] = len(outliers)

print(
    column,
    "-> Outliers:",
    len(outliers)
)
```

```
===== OUTLIER DETECTION =====
Age -> Outliers: 0
Attrition -> Outliers: 237
BusinessTravel -> Outliers: 0
DailyRate -> Outliers: 0
Department -> Outliers: 0
DistanceFromHome -> Outliers: 0
Education -> Outliers: 0
EducationField -> Outliers: 0
EmployeeCount -> Outliers: 0
EmployeeNumber -> Outliers: 0
EnvironmentSatisfaction -> Outliers: 0
Gender -> Outliers: 0
HourlyRate -> Outliers: 0
JobInvolvement -> Outliers: 0
JobLevel -> Outliers: 0
JobRole -> Outliers: 0
JobSatisfaction -> Outliers: 0
MaritalStatus -> Outliers: 0
MonthlyIncome -> Outliers: 114
MonthlyRate -> Outliers: 0
NumCompaniesWorked -> Outliers: 52
Over18 -> Outliers: 0
OverTime -> Outliers: 0
PercentSalaryHike -> Outliers: 0
PerformanceRating -> Outliers: 226
RelationshipSatisfaction -> Outliers: 0
StandardHours -> Outliers: 0
StockOptionLevel -> Outliers: 85
TotalWorkingYears -> Outliers: 63
TrainingTimesLastYear -> Outliers: 238
WorkLifeBalance -> Outliers: 0
YearsAtCompany -> Outliers: 104
YearsInCurrentRole -> Outliers: 21
YearsSinceLastPromotion -> Outliers: 107
YearsWithCurrManager -> Outliers: 14
```

Step 13: Visualize Outliers using Box Plots

```

boxplot_columns = encoded_df.select_dtypes(
    include=np.number
).columns

if len(boxplot_columns) > 0:

    plt.figure(figsize=(12, 6))


    sns.boxplot(
        data=encoded_df[boxplot_columns]
    )

    plt.title("Box Plot for Outlier Detection")
    plt.xticks(rotation=45)
    plt.tight_layout()

    plt.show()

print("Box Plot generated successfully.")

```



Step 14: Treat Outliers using IQR Capping

```

for column in numeric_columns:

    Q1 = encoded_df[column].quantile(0.25)
    Q3 = encoded_df[column].quantile(0.75)

    IQR = Q3 - Q1

    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    encoded_df[column] = encoded_df[column].clip(
        lower_bound,
        upper_bound
    )

print("Outliers treated using IQR Capping.")

Outliers treated using IQR Capping.

```

Step 15: Normalize / Standardize Numerical Attributes

```

numeric_columns = encoded_df.select_dtypes(
    include=np.number
).columns

scaler = StandardScaler()

if len(numeric_columns) > 0:

    encoded_df[numeric_columns] = scaler.fit_transform(
        encoded_df[numeric_columns]
    )

print("\n===== FEATURE SCALING =====")

print("Numerical attributes standardized successfully.")
print("Scaling Method: Standardization")
print("Mean approximately equals 0.")
print("Standard deviation approximately equals 1.")

===== FEATURE SCALING =====
Numerical attributes standardized successfully.
Scaling Method: Standardization
Mean approximately equals 0.
Standard deviation approximately equals 1.

```

Step 16: Feature Engineering

```

if "sibsp" in encoded_df.columns and "parch" in encoded_df.columns:

    encoded_df["family_size"] = (
        encoded_df["sibsp"] +
        encoded_df["parch"] +
        1
    )

```

```

)

print("\nFeature Created: family_size")

# If Age is available, create AgeGroup
if "age" in df.columns:

    age_values = df["age"]

    encoded_df["age_group"] = pd.cut(
        age_values,
        bins=[0, 12, 18, 35, 60, 100],
        labels=[
            "child",
            "teenager",
            "adult",
            "middle_aged",
            "senior"
        ]
    )

    # Encode Age Group
    encoded_df["age_group"] = LabelEncoder().fit_transform(
        encoded_df["age_group"].astype(str)
    )

    print("Feature Created: age_group")

```

Step 17: Display Preprocessed Dataset

```

print("\n===== PREPROCESSED DATASET =====")

print(encoded_df.head())

```

```

===== PREPROCESSED DATASET =====
   Age  Attrition  BusinessTravel  DailyRate  Department \
0  0.446350      0.0      0.590048   0.742527   1.401512
1  1.322365      0.0     -0.913194  -1.297775   -0.493817
2  0.008343      0.0      0.590048   1.414363   -0.493817
3 -0.429664      0.0     -0.913194   1.461466   -0.493817
4 -1.086676      0.0      0.590048  -0.524295   -0.493817

   DistanceFromHome  Education  EducationField  EmployeeCount  EmployeeNumber \
0      -1.010909    -0.891688    -0.937414         0.0      -1.701283
1      -0.147150   -1.868426    -0.937414         0.0     -1.699621
2      -0.887515   -0.891688     1.316673         0.0     -1.696298
3      -0.764121   1.061787    -0.937414         0.0     -1.694636
4      -0.887515   -1.868426     0.565311         0.0     -1.691313

   ...  RelationshipSatisfaction  StandardHours  StockOptionLevel \
0  ...      -1.584178         0.0      -0.978160
1  ...      1.191438         0.0      0.300537
2  ...     -0.658973         0.0     -0.978160
3  ...      0.266233         0.0     -0.978160
4  ...      1.191438         0.0      0.300537

   TotalWorkingYears  TrainingTimesLastYear  WorkLifeBalance  YearsAtCompany \
0      -0.424999      -2.130765      -2.493820     -0.123832
1      -0.151076      0.278639      0.338096     0.678078
2      -0.561960      0.278639      0.338096    -1.326697
3      -0.424999      0.278639      0.338096     0.277123
4      -0.698921      0.278639      0.338096    -0.925742

   YearsInCurrentRole  YearsSinceLastPromotion  YearsWithCurrManager
0      -0.058474      -0.776139         0.253823
1      0.785616      -0.372486         0.822605
2     -1.183926     -0.776139     -1.168130
3      0.785616      0.434819     -1.168130
4     -0.621200      0.031166     -0.599348

[5 rows x 35 columns]

```

Step 18: Display Final Dataset Information

```

print("\n===== FINAL DATASET INFORMATION =====")

print("Rows:", encoded_df.shape[0])
print("Columns:", encoded_df.shape[1])

print("\nMissing Values:")
print(encoded_df.isnull().sum().sum())

print("\nDuplicate Records:")
print(encoded_df.duplicated().sum())

```

```

===== FINAL DATASET INFORMATION =====
Rows: 1470
Columns: 35

Missing Values:
0

Duplicate Records:
0

```

Step 19: Save Cleaned Dataset

```

output_file = "cleaned_preprocessed_dataset.csv"

encoded_df.to_csv(
    output_file,
    index=False
)

print("\n===== DATASET STORAGE =====")

print(
    "Cleaned Dataset Saved Successfully:",
    output_file
)

```

)

```
===== DATASET STORAGE =====
Cleaned Dataset Saved Successfully: cleaned_preprocessed_dataset.csv
```

Step 20: Final Output

```
print("\n===== DATA PREPROCESSING =====")

print("Dataset Collection: Completed")
print("Dataset Loading: Completed")
print("Data Inspection: Completed")
print("Missing Value Handling: Completed")
print("Duplicate Removal: Completed")
print("Categorical Standardization: Completed")
print("Categorical Encoding: Completed")
print("Outlier Detection: Completed")
print("Outlier Treatment: Completed")
print("Feature Standardization: Completed")
print("Feature Engineering: Completed")
print("Cleaned Dataset Export: Completed")

print("\nObservations:")

print("Missing numerical values were handled using the median.")
print("Missing categorical values were handled using the mode.")
print("Duplicate records were removed.")
print("Categorical values were standardized and encoded.")
print("Outliers were detected using the IQR method.")
print("Outliers were treated using IQR capping.")
print("Numerical features were standardized using StandardScaler.")
print("New features were created where applicable.")

print("\nExperiment Completed Successfully.")
```

```
===== DATA PREPROCESSING =====
Dataset Collection: Completed
Dataset Loading: Completed
Data Inspection: Completed
Missing Value Handling: Completed
Duplicate Removal: Completed
Categorical Standardization: Completed
Categorical Encoding: Completed
Outlier Detection: Completed
Outlier Treatment: Completed
Feature Standardization: Completed
Feature Engineering: Completed
Cleaned Dataset Export: Completed

Observations:
Missing numerical values were handled using the median.
Missing categorical values were handled using the mode.
Duplicate records were removed.
Categorical values were standardized and encoded.
Outliers were detected using the IQR method.
Outliers were treated using IQR capping.
Numerical features were standardized using StandardScaler.
New features were created where applicable.

Experiment Completed Successfully.
```

EXPERIMENT 1 - ANSWERS TO QUESTIONS

QUESTION 1: Why is data preprocessing considered one of the most important phases in data analytics?

Data preprocessing is important because real-world data is often incomplete, inconsistent, duplicated and noisy.

Before performing visualization, statistical analysis or machine learning, the data must be cleaned and prepared.

Data preprocessing helps to:

Handle missing values. Remove duplicate records. Correct inconsistent data. Detect and treat outliers. Convert categorical variables into numerical form. Bring numerical features to a common scale.

Proper preprocessing improves the quality and reliability of analytical results.

QUESTION 2: Explain different methods of handling missing values with suitable examples.

Missing values can be handled using several methods.

Mean: Missing numerical values can be replaced by the mean of the available values.

Median: Missing numerical values can be replaced by the median. Median is useful when the data contains outliers.

Mode: Missing categorical values can be replaced by the most frequently occurring value.

Row Removal: Rows containing missing values can be removed when the number of missing records is small.

Column Removal: A column can be removed when a very large portion of its values is missing.

Example:

If the Age column contains missing values:

```
df["Age"] = df["Age"].fillna(df["Age"].median())
```

This replaces missing Age values with the median age.

QUESTION 3: Differentiate between Label Encoding and One-Hot Encoding.

Label Encoding:

Label Encoding converts categories into numerical labels.

Example:

Male → 0 Female → 1

It is useful when categories have an inherent order or when a compact numerical representation is required.

One-Hot Encoding:

One-Hot Encoding creates separate binary columns for each category.

Example:

Gender

Male Female

becomes:

Male Female 1 0 0 1

Main Difference:

Label Encoding → One numerical value per category. One-Hot Encoding → Separate binary column for each category.

QUESTION 4: What are outliers? How can they affect analytical results?

Outliers are observations that are significantly different from the majority of values in a dataset.

They can occur because of:

Measurement errors. Data entry errors. Unusual events. Natural variation.

Outliers can affect:

Mean and standard deviation. Statistical analysis. Correlation. Machine learning models. Data visualization.

The IQR method can be used to identify outliers.

$IQR = Q3 - Q1$

$Lower\ Bound = Q1 - 1.5 \times IQR$

$Upper\ Bound = Q3 + 1.5 \times IQR$

QUESTION 5: Explain the difference between normalization and standardization.

Normalization:

Normalization generally scales values to a fixed range, commonly 0 to 1.

A common formula is:

$X_{normalized} = (X - X_{min}) / (X_{max} - X_{min})$

Standardization:

Standardization transforms values so that the feature has approximately mean 0 and standard deviation 1.

Formula:

$X_{standardized} = (X - Mean) / Standard\ Deviation$

Main Difference:

Normalization → Fixed range, commonly 0 to 1. Standardization → Mean approximately 0 and standard deviation approximately 1.

QUESTION 6: Why should duplicate records be removed before analysis?

Duplicate records represent repeated observations of the same data.

If duplicates are not removed, they can:

Increase the apparent frequency of certain observations. Produce biased statistical results. Affect averages and distributions. Affect machine learning model training. Increase the apparent size of the dataset.

Removing duplicate records helps ensure that each observation is represented appropriately in the dataset.

QUESTION 7: What is feature engineering? Give two practical examples.

Feature engineering is the process of creating new useful attributes from existing data.

It can help represent information in a form that is more useful for analysis or machine learning.

Example 1:

Family Size can be created from:

$SibSp + Parch + 1$

$FamilySize = SibSp + Parch + 1$

Example 2:

Age Group can be created from Age.

For example:

0-12 → Child 13-18 → Teenager 19-35 → Adult 36-60 → Middle Aged 60+ → Senior

Feature engineering can provide additional useful information without collecting new raw data.

QUESTION 8: Which preprocessing techniques would you apply to the IBM HR Employee Attrition dataset and why?

For the IBM HR Employee Attrition dataset, the following preprocessing techniques can be applied:

Missing Value Handling: Identify and appropriately handle missing values.

Duplicate Removal: Remove duplicate employee records.

Categorical Encoding: Convert categorical attributes such as department, job role and marital status into numerical form.

Outlier Detection: Use box plots and the IQR method for numerical variables.

Feature Scaling: Standardize numerical features such as age, income and years of experience where appropriate.

Feature Engineering: Create useful attributes such as age groups or income categories.

These techniques improve data quality and prepare the employee dataset for statistical analysis and machine learning.

QUESTION 9: How does poor-quality data affect machine learning model performance?

Poor-quality data can negatively affect machine learning models.

Problems such as:

Missing values. Duplicate records. Incorrect values. Outliers. Inconsistent categories. Different numerical scales.

can cause models to learn incorrect or misleading patterns.

This may result in:

Lower prediction accuracy. Poor generalization. Biased predictions. Unstable models. Reduced reliability.

Therefore, high-quality preprocessing is important for building reliable machine learning models.

QUESTION 10: Name any three Python libraries commonly used for data preprocessing.

Three commonly used Python libraries are:

Pandas: Used for loading, cleaning, manipulating and analyzing structured datasets.

NumPy: Used for numerical operations and array processing.

Scikit-learn: Provides preprocessing tools such as LabelEncoder, StandardScaler and other machine learning utilities.

Other useful libraries include Matplotlib and Seaborn for visualization and exploratory data analysis.

