

▼ Name = Dev Sharma

CU24250269
BTech CSE 3rd Year SECTION - "B"
Roll No. :- 17

Experiment Name
Exploratory Data Analysis (EDA) on a Real-World Business
Dataset using Python

Aim:
To perform Exploratory Data Analysis (EDA) on a real-world dataset using Python in order to understand the dataset's structure, identify trends, detect anomalies, analyze feature relationships, and generate meaningful business insights through descriptive statistics and visualizations.

GITHUB LINK :- [https://github.com/DevSharma1912/Mastery-in-Data-Analytics-and-Visualizations-and-Career-Advancement-Technical-Tree/main/Mastery%20in%20Data%20Analytics%20and%20Visualizations%20and%20Career%20Advancement%20\(Technical\)/Experiment2](https://github.com/DevSharma1912/Mastery-in-Data-Analytics-and-Visualizations-and-Career-Advancement-Technical-Tree/main/Mastery%20in%20Data%20Analytics%20and%20Visualizations%20and%20Career%20Advancement%20(Technical)/Experiment2)

Step 1: Import Required Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from tkinter import Tk, filedialog
```

step 2: Select / Upload Dataset from Computer

```
from google.colab import files

print("Please upload your cleaned business dataset.")
uploaded = files.upload()

if not uploaded:
    print("No dataset selected.")
    file_path = None
else:
    # Get the name of the first uploaded file
    file_path = next(iter(uploaded.keys()))
    print(f"File '{file_path}' uploaded successfully.")
```

Please upload your cleaned business dataset.
[Choose File](#) No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving Mall_Customers.csv to Mall_Customers.csv
File 'Mall_Customers.csv' uploaded successfully.

Step 3: Load Dataset

```
if file_path.lower().endswith(".csv"):
    df = pd.read_csv(file_path)

elif file_path.lower().endswith((".xlsx", ".xls")):
    df = pd.read_excel(file_path)

else:
    print("Unsupported file format.")
    # Instead of exiting, we'll set df to None or handle as appropriate
    df = None # Or raise an error if strict file type is required

if df is not None:
    print("Dataset loaded successfully.")
    print("Selected Dataset:", file_path)
else:
    print("Failed to load dataset.")
```

Dataset loaded successfully.
Selected Dataset: Mall_Customers.csv

Step 4: Display First Five Records

```
print("\n===== FIRST FIVE RECORDS =====")

print(df.head())
```

===== FIRST FIVE RECORDS =====					
	CustomerID	Genre	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

Step 5: Display Dataset Dimensions

```
print("\n===== DATASET DIMENSIONS =====")

print("Number of Rows:", df.shape[0])
print("Number of Columns:", df.shape[1])
```

===== DATASET DIMENSIONS =====
Number of Rows: 200
Number of Columns: 5

step 6: Display Column Names

```
print("\n===== COLUMN NAMES =====")

print(df.columns.tolist())
```

===== COLUMN NAMES =====
['CustomerID', 'Genre', 'Age', 'Annual Income (k\$)', 'Spending Score (1-100)']

Step 7: Display Data Types

```
print("\n===== DATA TYPES =====")

print(df.dtypes)
```

===== DATA TYPES =====
CustomerID int64

```
Genre          object
Age            int64
Annual Income (k$)  int64
Spending Score (1-100)  int64
dtype: object
```

Step 8: Display Dataset Information

```
print("\n===== DATASET INFORMATION =====")

df.info()

===== DATASET INFORMATION =====
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column              Non-Null Count  Dtype  
---  --
0   CustomerID          200 non-null   int64  
1   Genre               200 non-null   object  
2   Age                 200 non-null   int64  
3   Annual Income (k$)  200 non-null   int64  
4   Spending Score (1-100)  200 non-null   int64  
dtypes: int64(4), object(1)
memory usage: 7.5+ KB
```

Step 9: Statistical Summary

```
print("\n===== STATISTICAL SUMMARY =====")

print(df.describe(include="all"))

===== STATISTICAL SUMMARY =====
count      CustomerID      Genre      Age      Annual Income (k$) \
unique           NaN          2           NaN          NaN
top              NaN      Female           NaN          NaN
freq           NaN          112           NaN          NaN
mean         100.500000      NaN      38.850000      60.560000
std           57.873115      NaN      11.969007      26.264721
min            1.000000      NaN      18.000000      15.000000
25%           50.750000      NaN      28.750000      41.500000
50%          100.500000      NaN      36.000000      61.500000
75%          150.250000      NaN      49.000000      78.000000
max           200.000000      NaN      70.000000      137.000000

      Spending Score (1-100)
count           200.000000
unique           NaN
top              NaN
freq           NaN
mean           50.280000
std            25.823522
min             1.000000
25%            34.750000
50%            50.000000
75%            73.000000
max            99.000000
```

Step 10: Missing Value Analysis

```
print("\n===== MISSING VALUES =====")

missing_values = df.isnull().sum()

print(missing_values)

print("\nTotal Missing Values:",
      df.isnull().sum().sum())

===== MISSING VALUES =====
CustomerID      0
Genre           0
Age             0
Annual Income (k$)  0
Spending Score (1-100)  0
dtype: int64

Total Missing Values: 0
```

Step 11: Duplicate Analysis

```
print("\n===== DUPLICATE RECORDS =====")

duplicate_count = df.duplicated().sum()

print("Number of Duplicate Records:",
      duplicate_count)

===== DUPLICATE RECORDS =====
Number of Duplicate Records: 0
```

Step 12: Identify Numerical and Categorical Columns

```
numerical_columns = df.select_dtypes(
    include=np.number
).columns.tolist()

categorical_columns = df.select_dtypes(
    exclude=np.number
).columns.tolist()

print("\n===== NUMERICAL FEATURES =====")

print(numerical_columns)

print("\n===== CATEGORICAL FEATURES =====")

print(categorical_columns)

===== NUMERICAL FEATURES =====
['CustomerID', 'Age', 'Annual Income (k$)', 'Spending Score (1-100)']

===== CATEGORICAL FEATURES =====
['Genre']
```

Step 13: Value Counts for Categorical Variables

```
print("\n===== CATEGORICAL VALUE COUNTS =====")

for column in categorical_columns:

    print("\nColumn:", column)

    print(
        df[column]
        .value_counts()
        .head(10)
    )

===== CATEGORICAL VALUE COUNTS =====
Column: Genre
Genre
Female    112
Male      88
Name: count, dtype: int64
```

Step 14: Univariate Analysis - Histograms

```
print("\n===== UNIVARIATE ANALYSIS =====")

if len(numerical_columns) > 0:

    for column in numerical_columns:

        plt.figure(figsize=(8, 5))

        sns.histplot(
            data=df,
            x=column,
            kde=True
        )


        plt.title(
            "Distribution of " + column
        )

        plt.xlabel(column)
        plt.ylabel("Frequency")

        plt.tight_layout()
        plt.show()

print("Numerical feature distributions analyzed.")
```

===== UNIVARIATE ANALYSIS =====



Step 15: Categorical Variable Analysis

```
if len(categorical_columns) > 0:

    # Select first few categorical columns
    selected_categories = categorical_columns[:4]

    for column in selected_categories:

        # Limit categories for better visualization
        top_categories = (
            df[column]
            .value_counts()
            .head(10)
        )

        plt.figure(figsize=(10, 5))
```


```
sns.barplot(
    x=top_categories.values,
    y=top_categories.index
)

plt.title(
    "Top Categories in " + column
)

plt.xlabel("Count")
plt.ylabel(column)

plt.tight_layout()
plt.show()

print("Categorical variables analyzed using bar charts.")
```



Step 16: Box Plot for Numerical Features
Categorical variables analyzed using bar charts.

```
if len(numerical_columns) > 0:

    plt.figure(figsize=(12, 6))


    sns.boxplot(
        data=df[numerical_columns]
    )

    plt.title(
        "Box Plot of Numerical Features"
    )

    plt.xticks(rotation=45)

    plt.tight_layout()
    plt.show()

print("Box plot generated successfully.")
```



Box plot generated successfully.

Step 17: Correlation Matrix

```
if len(numerical_columns) >= 2:

    correlation_matrix = df[
        numerical_columns
    ].corr()

    print("\n===== CORRELATION MATRIX =====")

    print(correlation_matrix)
```

===== CORRELATION MATRIX =====			
	CustomerID	Age	Annual Income (k\$) \
CustomerID	1.000000	-0.026763	0.977548
Age	-0.026763	1.000000	-0.012398
Annual Income (k\$)	0.977548	-0.012398	1.000000
Spending Score (1-100)	0.013835	-0.327227	0.009983

	Spending Score (1-100)
CustomerID	0.013835
Age	-0.327227
Annual Income (k\$)	0.009983
Spending Score (1-100)	1.000000

Step 18: Correlation Heatmap

```
if len(numerical_columns) >= 2:

    plt.figure(
        figsize=(12, 8)
    )

    sns.heatmap(
        correlation_matrix,
        annot=True,
        cmap="coolwarm",
        fmt=".2f",
        linewidths=0.5
    )


    plt.title(
        "Correlation Heatmap"
    )

    plt.tight_layout()
```

```
plt.show()

print(
    "Correlation heatmap generated successfully."
)

else:
    print(
        "\nNot enough numerical features "
        "for correlation analysis."
    )
```



Step 19: Bivariate Analysis

```
print("\n===== BIVARIATE ANALYSIS =====")

if len(numerical_columns) >= 2:
    x_column = numerical_columns[0]
    y_column = numerical_columns[1]

    plt.figure(figsize=(8, 6))

    sns.scatterplot(
        data=df,
        x=x_column,
        y=y_column
    )


    plt.title(
        x_column + " vs " + y_column
    )

    plt.xlabel(x_column)
    plt.ylabel(y_column)

    plt.tight_layout()
    plt.show()

    print(
        "Scatter plot generated for:",
        x_column,
        "and",
        y_column
    )

else:
    print(
        "Not enough numerical columns "
        "for scatter plot."
    )
```



Step 20: Multivariate Analysis

```
print("\n===== MULTIVARIATE ANALYSIS =====")

if len(numerical_columns) >= 3:
    sns.pairplot(
        df[numerical_columns[:5]].dropna()
    )


    plt.show()

    print(
        "Pair plot generated successfully."
    )
```

```

)
else:
    print(
        "Not enough numerical features "
        "for multivariate analysis."
    )

```



Step 21: Mean, Median and Standard Deviation

```

print("\n===== DESCRIPTIVE STATISTICS =====")

if len(numerical_columns) > 0:
    for column in numerical_columns:
        print("\nFeature:", column)

        print(
            "Mean:",
            round(df[column].mean(), 2)
        )

        print(
            "Median:",
            round(df[column].median(), 2)
        )

        print(
            "Standard Deviation:",
            round(df[column].std(), 2)
        )

        print(
            "Minimum:",
            df[column].min()
        )

        print(
            "Maximum:",
            df[column].max()
        )

```

```

===== DESCRIPTIVE STATISTICS =====

Feature: CustomerID
Mean: 100.5
Median: 100.5
Standard Deviation: 57.88
Minimum: 1
Maximum: 200

Feature: Age
Mean: 38.85
Median: 36.0
Standard Deviation: 13.97
Minimum: 18
Maximum: 70

Feature: Annual Income (k$)
Mean: 60.56
Median: 63.5
Standard Deviation: 26.26
Minimum: 15
Maximum: 137

Feature: Spending Score (1-100)
Mean: 50.2
Median: 50.0
Standard Deviation: 25.82
Minimum: 1
Maximum: 99

```

Step 22: Detect Highly Correlated Features

```

if len(numerical_columns) >= 2:
    correlation_matrix = df[
        numerical_columns
    ].corr()

    print(
        "\n===== HIGH CORRELATIONS ====="
    )

    found_correlation = False

    for i in range(
        len(correlation_matrix.columns)
    ):

```

```
for j in range(i + 1,
               len(correlation_matrix.columns)):

    correlation_value = (
        correlation_matrix.iloc[i, j]
    )

    if abs(correlation_value) >= 0.70:

        print(
            correlation_matrix.columns[i],
            "<->",
            correlation_matrix.columns[j],
            ":",
            round(correlation_value, 2)
        )

        found_correlation = True

if not found_correlation:

    print(
        "No strong correlation "
        "(absolute correlation >= 0.70) found."
    )

===== HIGH CORRELATIONS =====
CustomerID <-> Annual Income (k$) : 0.98
```

Step 23: Business Insights

```
print("\n===== BUSINESS INSIGHTS =====")

print(
    "1. Numerical distributions were analyzed "
    "using histograms and density plots."
)

print(
    "2. Frequently occurring categories were "
    "identified using value counts and bar charts."
)

print(
    "3. Box plots were used to identify "
    "potential outliers."
)

print(
    "4. Correlation analysis was used to "
    "identify relationships between numerical features."
)

print(
    "5. Scatter plots were used to analyze "
    "relationships between pairs of variables."
)

print(
    "6. Multivariate analysis was used to "
    "observe relationships among multiple features."
)

===== BUSINESS INSIGHTS =====
1. Numerical distributions were analyzed using histograms and density plots.
2. Frequently occurring categories were identified using value counts and bar charts.
3. Box plots were used to identify potential outliers.
4. Correlation analysis was used to identify relationships between numerical features.
5. Scatter plots were used to analyze relationships between pairs of variables.
6. Multivariate analysis was used to observe relationships among multiple features.
```

Step 24: Final Output

```
print(
    "\n===== EXPLORATORY DATA ANALYSIS ====="
)

print("Dataset Loading: Completed")
print("Dataset Inspection: Completed")
print("Statistical Analysis: Completed")
print("Missing Value Analysis: Completed")
print("Duplicate Analysis: Completed")
print("Univariate Analysis: Completed")
print("Categorical Analysis: Completed")
print("Correlation Analysis: Completed")
print("Heatmap Generation: Completed")
print("Box Plot Analysis: Completed")
print("Scatter Plot Analysis: Completed")
print("Multivariate Analysis: Completed")
print("Business Insight Generation: Completed")

print("\nExperiment Completed Successfully.")

===== EXPLORATORY DATA ANALYSIS =====
Dataset Loading: Completed
Dataset Inspection: Completed
Statistical Analysis: Completed
Missing Value Analysis: Completed
Duplicate Analysis: Completed
Univariate Analysis: Completed
Categorical Analysis: Completed
Correlation Analysis: Completed
Heatmap Generation: Completed
Box Plot Analysis: Completed
Scatter Plot Analysis: Completed
Multivariate Analysis: Completed
Business Insight Generation: Completed

Experiment Completed Successfully.
```

EXPERIMENT 2 - ANSWERS TO QUESTIONS

QUESTION 1: What is Exploratory Data Analysis (EDA), and why is it performed before machine learning?

Exploratory Data Analysis (EDA) is the process of examining, summarizing and visualizing a dataset to understand its structure, patterns, relationships and anomalies.

EDA is performed before machine learning because it helps identify:

Missing values. Duplicate records. Outliers. Data distributions. Relationships between features. Skewness. Class imbalance.

EDA helps analysts understand the data before selecting features, preprocessing techniques and machine learning models.

QUESTION 2: Differentiate between univariate, bivariate, and multivariate analysis with suitable examples.

Univariate Analysis:

Analysis of a single variable.

Example: Studying the distribution of employee age using a histogram.

Bivariate Analysis:

Analysis of two variables to understand their relationship.

Example: Studying the relationship between sales and profit using a scatter plot.

Multivariate Analysis:

Analysis involving three or more variables simultaneously.
Example: Studying relationships between sales, profit, quantity and discount using a pair plot or multivariate visualization.
Therefore:

Univariate → One variable. Bivariate → Two variables. Multivariate → Three or more variables.

QUESTION 3: What insights can be obtained from a correlation heatmap?

A correlation heatmap represents the correlation between numerical variables using colors and correlation values.

It can help identify:

Strong positive relationships. Strong negative relationships. Weak relationships. Highly correlated features. Possible multicollinearity.

Correlation values generally range from -1 to +1.

+1 → Strong positive correlation. -1 → Strong negative correlation. 0 → Little or no linear correlation.

A heatmap therefore provides a quick visual summary of relationships between numerical features.

QUESTION 4: Explain the purpose of histograms, box plots, and scatter plots in EDA.

Histogram:

A histogram shows the distribution of a numerical variable. It helps identify:

Central tendency. Spread. Skewness. Possible multiple peaks.

Box Plot:

A box plot summarizes the distribution of a numerical variable and helps identify possible outliers.

It displays:

Median. Quartiles. Interquartile range. Potential outliers.

Scatter Plot:

A scatter plot shows the relationship between two numerical variables.

It helps identify:

Positive relationships. Negative relationships. Clusters. Outliers. Patterns.

QUESTION 5: How can EDA help identify data quality issues before analysis?

EDA helps identify data quality problems by examining the structure, values and distributions of the dataset.

It can detect:

Missing values using isnull(). Duplicate records using duplicated(). Incorrect data types using info() and dtypes. Outliers using box plots.

Inconsistent categories using value_counts(). Unusual distributions using histograms. Extreme values using descriptive statistics.

Identifying these issues before analysis improves the reliability of the results.

QUESTION 6: Why is correlation important in predictive analytics? Can correlation imply causation?

Correlation is important because it measures the strength and direction of a linear relationship between variables.

In predictive analytics, correlation can help identify potentially useful features and detect highly related variables.

However, correlation does NOT necessarily imply causation.

Two variables can be correlated because:

One variable influences another. Both are influenced by another variable. The relationship occurs by coincidence.

Therefore, correlation indicates association, not necessarily a cause-and-effect relationship.

QUESTION 7: Which visualization would you use to analyze categorical and numerical variables? Justify your choice.

For categorical variables:

Bar charts and count plots are suitable because they show the frequency of different categories clearly.

For numerical variables:

Histograms are useful for understanding the distribution.

Box plots are useful for comparing distributions and identifying outliers.

For relationships between numerical variables:

Scatter plots are useful because they display the relationship between two numerical variables.

Therefore:

Categorical → Bar Chart / Count Plot. Numerical → Histogram / Box Plot. Two Numerical Variables → Scatter Plot.

QUESTION 8: What business insights can be derived from the Netflix (or Superstore/HR Analytics) dataset through EDA?

EDA can provide different business insights depending on the selected dataset.

Netflix Dataset:

Distribution of movies and TV shows. Most common genres. Content distribution by country. Content added over time. Distribution of movie durations.

Superstore Dataset:

Sales performance. Profitability. Product category performance. Regional performance. Relationship between discount and profit.

IBM HR Dataset:

Employee attrition patterns. Job-role distribution. Salary-related patterns. Department-wise employee characteristics. Factors associated with employee attrition.

Therefore, EDA can help organizations identify trends, anomalies and factors that may influence business decisions.

QUESTION 9: How does EDA contribute to feature selection and model building?

EDA helps understand which features may contain useful information for predicting the target variable.

Correlation analysis can identify relationships between numerical features.

Distribution analysis can identify highly skewed variables.

Box plots can identify extreme values.

Categorical analysis can reveal important groups or classes.

EDA can also identify redundant or highly correlated features that may create multicollinearity.

Therefore, EDA helps in:

Selecting useful features. Removing irrelevant features. Detecting redundant features. Choosing suitable preprocessing techniques.

Selecting appropriate machine learning models.

QUESTION 10: What challenges might arise while performing EDA on large-scale real-world datasets?

EDA on large datasets can present several challenges.

High computational requirements: Large datasets require more CPU and memory.

Large storage requirements: Very large datasets may require substantial storage.

Slow visualization: Plotting millions of records can be time-consuming.

Missing and inconsistent data: Large datasets may contain many data quality problems.

High dimensionality: Datasets with many features can be difficult to visualize and analyze.

Processing time: Statistical calculations can take longer on very large datasets.

Sampling, efficient data structures, aggregation and distributed processing can be used to make EDA more manageable on large-scale datasets.